



University of Reading  
Department of Computer Science

## CS3AI18: Predicting Rainfall in Australia

Savia Laloukiotou

Student Number: 28017133  
Module Convenor: Dr Yevgeniya Kovalchuk

March 7, 2022

# Contents

<b>1</b>	<b>Background</b>	<b>2</b>
<b>2</b>	<b>Dataset</b>	<b>3</b>
<b>3</b>	<b>Machine Learning Model</b>	<b>4</b>
3.1	Solution approach . . . . .	4
3.2	Data Preprocessing and Visualisation . . . . .	4
3.3	Model Training and Evaluation . . . . .	12
3.4	Results . . . . .	18
<b>4</b>	<b>Conclusions and Future Work</b>	<b>21</b>

## **Abstract**

Rainfall forecast is a challenging endeavour. It plays a critical role in decision making and the organisation of sustainable water resource systems. Although meteorological systems produce short-term rainfall forecasts, long-term rainfall forecasting is difficult and fraught with uncertainty due to a variety of causes. Finding the optimum approach to predict the complex rainfall process in areas like Australia, where the climate is erratic, is a major challenge. The aim of this project is to build a machine learning model to address the problem of rainfall prediction. The results demonstrate that the Random Forest algorithm outperforms Logistic Regression in predicting rainfall.

# Chapter 1

## Background

Rainfall prediction remains a major concern, attracting the attention of governments, businesses, risk management organisations, and researchers. Rainfall is a climatic component that has an impact on a variety of human activities, including agriculture, building, power generation, forestry, and tourism (WHO, 2003). Its forecasting is crucial in as it is closely linked to severe natural events including landslides, flooding, mass movements, and avalanches. The continent is subjected to a wide range of rainfall patterns, including floods and droughts (Sharmila, 2020). Developing a successful rainfall prediction model allows to take preventive and mitigating measures against these natural occurrences (Nicholls, 2001).

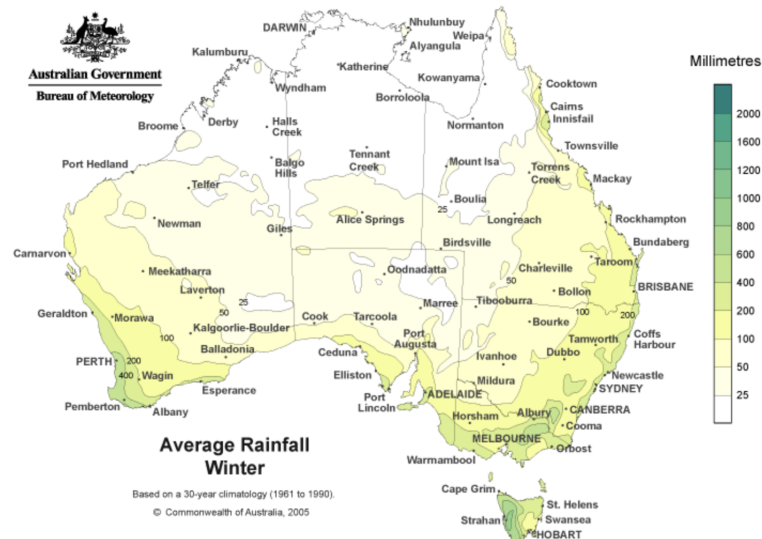


Figure 1.1: Average Winter Rainfall in Australia based on 30 year climatology.

## Chapter 2

# Dataset

The dataset used contains 10 years' worth of daily weather observations from multiple Australian weather stations (*Australian Government Bureau of Meteorology*, n.d.) and consists of 145460 records and 23 attributes. Table 2.1 shows a description of the attributes included. RainTomorrow is the target variable.

Name	Description
Date	The date of observation
Location	Location name of the weather station
MinTemp	The minimum temperature in degrees celsius
MaxTemp	The maximum temperature in degrees celsius
Rainfall	The amount of rainfall recorded for the day in mm
Evaporation	The Class A pan evaporation (mm) in the 24 hours to 9am
Sunshine	The number of hours of bright sunshine in the day
WindGustDir	The direction of the strongest wind gust in the 24 hours to midnight
WindGustSpeed	The speed (km/h) of the strongest wind gust in the 24 hours to midnight
WindDir9am	Direction of the wind at 9am
WindDir3pm	Direction of the wind at 3pm
WindSpeed9am	Wind speed (km/hr) averaged over 10 minutes prior to 9am
WindSpeed3pm	Wind speed (km/hr) averaged over 10 minutes prior to 3pm
Humidity9am	Humidity (percent) at 9am
Humidity3pm	Humidity (percent) at 3pm
Pressure9am	Atmospheric pressure (hpa) reduced to mean sea level at 9am
Pressure3pm	Atmospheric pressure (hpa) reduced to mean sea level at 3pm
Cloud9am	Fraction of sky obscured by cloud (in "oktas": eighths) at 9am
Cloud3pm	Fraction of sky obscured by cloud (in "oktas": eighths) at 3pm
Temp9am	Temperature (degrees C) at 9am
Temp3pm	Temperature (degrees C) at 3pm
RainToday	"Yes" if precipitation (mm) in the 24 hours to 9am exceeds 1mm, otherwise "No"
RainTomorrow	"Yes" if precipitation (mm) in the next day exceeds 1mm, otherwise "No"

Table 2.1: Names and description of the attributes.

## Chapter 3

# Machine Learning Model

### 3.1 Solution approach

Sharma et al. (2021) used the Decision Tree algorithm to construct a rainfall forecast model with an accuracy of 81%. Gupta and Ghose (2015) used the KNN (K NearestNeighbors) and Naive Bayes algorithms, with 80.7% and 78.9% accuracy respectively. However, none of the aforementioned papers mention the F1 score or other metrics. As described in the following chapters, the accuracy of a model is not always a good measure of its performance. In this approach, Random Forest and Logistic Regression algorithms are going to be implemented to test whether they can achieve improved results.

### 3.2 Data Preprocessing and Visualisation

Using the `info()` and `describe()` methods (Listing 3.1) we can retrieve useful insights about our data. The `info()` method returns the non-null count and type of each variable (Figure 3.1) and `describe()` returns a statistical summary of the numerical variables (Figure 3.2).

```
1  ### VIEW THE SHAPE AND HAVE A QUICK LOOK AT THE DATA
2
3  print(f'The number of rows are {df.shape[0]} and the number of columns
      are {df.shape[1]}')
4  pd.set_option('display.max_columns', None) # to view all the columns
5  print(df.head()) # to have a quick look at the data
6
7  ###
8  print(df.info()) # print a summary of the variables in the dataframe
9
10 ###
11 print(df.describe()) # view a statistics summary of the numeric data
```

Listing 3.1: Information and statistical summary of the dataset

We can see that the dataset contains both categorical and numerical variables. Categorical variables have data type object and numerical variables have data type float64. It also appears that some columns contain a high amount of missing values. From the statistical summary, we can see that some variables have a higher standard deviation value which indicates that their data points are dispersed. A large difference between the 75th percentile and the maximum value of some variables indicates the presence of potential outliers (e.g., the max Rainfall value is 371.0 despite the 75th percentile being 0.8). The same goes for the difference between the

Data #	Column	Non-Null Count	Dtype
0	Date	145460 non-null	object
1	Location	145460 non-null	object
2	MinTemp	143975 non-null	float64
3	MaxTemp	144199 non-null	float64
4	Rainfall	142199 non-null	float64
5	Evaporation	82670 non-null	float64
6	Sunshine	75625 non-null	float64
7	WindGustDir	135134 non-null	object
8	WindGustSpeed	135197 non-null	float64
9	WindDir9am	134894 non-null	object
10	WindDir3pm	141232 non-null	object
11	WindSpeed9am	143693 non-null	float64
12	WindSpeed3pm	142398 non-null	float64
13	Humidity9am	142806 non-null	float64
14	Humidity3pm	140953 non-null	float64
15	Pressure9am	130395 non-null	float64
16	Pressure3pm	130432 non-null	float64
17	Cloud9am	89572 non-null	float64
18	Cloud3pm	86102 non-null	float64
19	Temp9am	143693 non-null	float64
20	Temp3pm	141851 non-null	float64
21	RainToday	142199 non-null	object
22	RainTomorrow	142193 non-null	object

dtypes: float64(16), object(7)

Figure 3.1: Information about the attributes.

	MinTemp	MaxTemp	Rainfall	Evaporation
count	143975.000000	144199.000000	142199.000000	82670.000000
mean	12.194034	23.221348	2.360918	5.468232
std	6.398495	7.119049	4.800000	4.193704
min	-8.500000	-4.800000	0.000000	0.000000
25%	7.500000	17.900000	0.000000	2.600000
50%	12.000000	22.600000	0.000000	4.800000
75%	16.900000	28.200000	0.000000	7.400000
max	33.900000	48.100000	371.000000	145.000000

	Sunshine	WindGustSpeed	WindSpeed9am	WindSpeed3pm
count	75625.000000	135197.000000	143693.000000	142398.000000
mean	7.611178	40.035230	14.043426	18.662657
std	3.785483	13.607062	8.915375	8.809800
min	0.000000	6.000000	0.000000	0.000000
25%	4.000000	31.000000	7.000000	13.000000
50%	8.400000	39.000000	13.000000	19.000000
75%	10.600000	48.000000	19.000000	24.000000
max	14.500000	135.000000	130.000000	87.000000

	Humidity9am	Humidity3pm	Pressure9am	Pressure3pm
count	142806.000000	140953.000000	130395.000000	130432.000000
mean	68.880831	51.539116	1017.64994	1015.255889
std	19.020164	20.795902	7.10653	7.037414
min	0.000000	0.000000	980.50000	977.100000
25%	57.000000	37.000000	1012.90000	1010.40000
50%	70.000000	52.000000	1017.60000	1015.20000
75%	83.000000	66.000000	1022.40000	1020.00000
max	100.000000	100.000000	1041.00000	1039.60000

	Cloud9am	Cloud3pm	Temp9am	Temp3pm
count	89572.000000	86102.000000	143693.000000	141851.000000
mean	4.447461	4.509930	16.990631	21.68339
std	2.887159	2.720357	6.488753	6.93665
min	0.000000	0.000000	-7.200000	-5.40000
25%	1.000000	2.000000	12.300000	16.60000
50%	5.000000	5.000000	16.700000	21.10000
75%	7.000000	7.000000	21.600000	26.40000
max	9.000000	9.000000	40.200000	46.70000

Figure 3.2: Statistical description of the numerical attributes.

minimum value and the 25th percentile. To further investigate how the numerical variables are distributed, we plot their histograms (Figure 3.3). The histograms show us that some variables follow normal distribution, (e.g. Humidity3pm, MinTemp, Temp9am), while some are highly skewed (e.g. Evaporation and Rainfall).

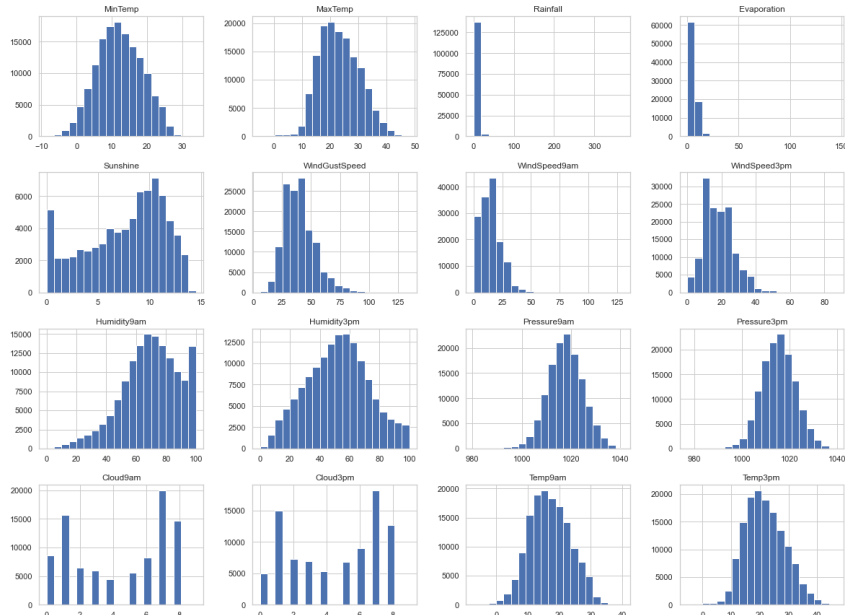


Figure 3.3: Distribution of the numerical variables.

We now investigate the correlation between the variables. In order to do so, we plot a correlation heatmap (Figure 3.4).




Figure 3.4: Variable correlation heatmap.

The feature pairs listed below are highly correlated:

- MaxTemp and MinTemp
- Pressure9am and Pressure3pm
- Temp9am and Temp3pm
- Evaporation and MaxTemp
- MaxTemp and Temp3pm

Strong correlated features (either positive or negative) have intuitive reasons for being so. However, by looking at the pair diagram below (Figure 3.5), we may learn more about the pairwise correlation between these features. The "yes" and "no" clusters of RainTomorrow are plainly visible. There is only a small amount of overlap.



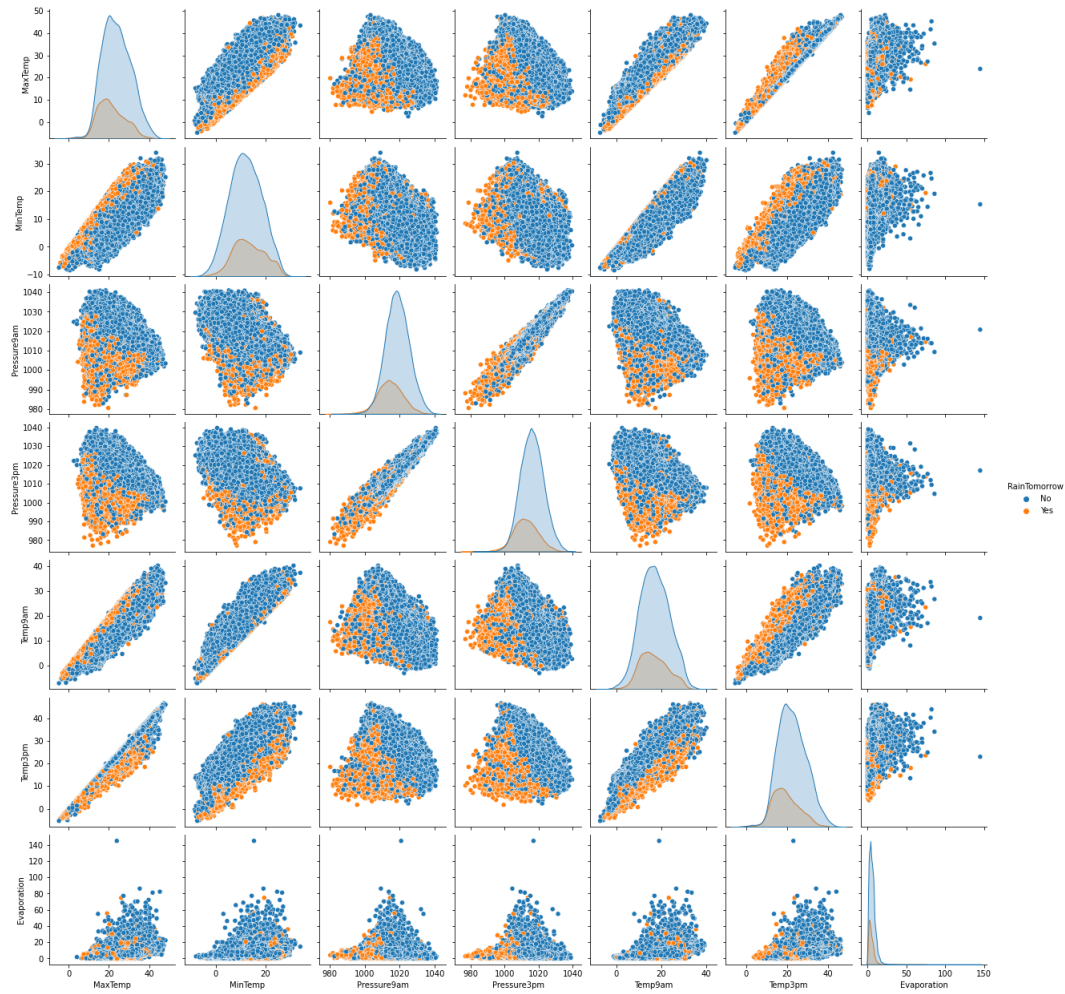


Figure 3.5: Pairplots of the variables.

Next, we have a closer look at the balance of the class variable `RainTomorrow`. As shown in Figure 3.6, the target class is imbalanced.

Balance of positive and negative classes:

0: 79.4%

1: 20.6%

Standard classifier algorithms have a bias towards classes which have large number of instances. The features of the minority class are treated as noise and are often ignored. Thus, there is a high probability of misclassification of the minority class as compared to the majority class (AnalyticsVidhya, 2017). In our case, if the classifier always predicts No rainfall, it will be correct around 80% of the time although in practice it is performing no better than an unskilled majority class classifier. This is also known as the Accuracy Paradox (MachineLearningMastery, 2021).




Figure 3.6: Target Class Distribution.

Then, we check for missing values in our dataset. When checking for missing values in categorical variables, checking for unique values and their count is essential to spot irrelevant entries as well (e.g. null values can be encoded as "?").

```
1 df_cat= df.select_dtypes(exclude=['float']) # include only object type
   variables
2 for col in df_cat.columns:
3     print(df_cat[col].unique()) # to print categories name only
4     print(df_cat[col].value_counts()) # to print count of every category
```

Listing 3.2: Printing unique values and value counts of categorical variables

The output of the code in Listing 3.2 showed that missing values are encoded as "nan". To visualize them better, we plot a heatmap (Figure 3.7).




Figure 3.7: Missing values heatmap.

It is clear that 'Evaporation', 'Sunshine', 'Cloud9am' and 'Cloud3pm' are the features that contain the most missing values. To handle missing values throughout all the columns, we first split the features to two lists, Numerical and Categorical. Then, using the `fillna()` method we replace the null values with the median value for numerical, and the mode value for categorical variables (Listing 3.3). The median is more robust (less sensitive to outliers in the data) than the mean (Wand, n.d.).

```

1 # List of categorical variables
2 categorical = [i for i in df.columns if df[i].dtypes == 'O']
3 # List of numerical variables
4 numerical = [i for i in df.columns if i not in categorical]
5 print('Categorical:\n', categorical, '\n\n', 'Numerical:\n', numerical)
6
7 %%
8 # function to replace null numerical values
9 def replace_numerical(df):
10     for col in numerical: # for all numerical columns
11         df[col] = df[col].fillna(df[col].median()) # replace missing value
12         with the median
13     return df
14
15 # function to replace null categorical values
16 def replace_categorical(df):
17     for col in categorical: # for all categorical columns
18         df[col] = df[col].fillna(df[col].mode()[0]) # replace missing
19         value with mode
20
21     return df
22
23 #%% REPLACE NUMERICAL AND CATEGORICAL VALUES
24 df = replace_numerical(df)
25 df = replace_categorical(df)

```

Listing 3.3: Replacing missing values

After replacing the missing values, we plot the missing values heatmap again (Figure 3.8) and observe that all null values have been eliminated.




Figure 3.8: Missing values heatmap after replacement.

Machine learning models require all input and output variables to be numeric. This means that categorical variables must be encoded. The "YES" and "NO" variables (i.e. RainToday and RainTomorrow) are encoded with 1 and 0. (Listing 3.4).

```

1  ### REPLACE YES OR NO VALUES WITH 1 AND 0
2  df = df.replace({'RainTomorrow': {'Yes': 1, 'No': 0}})
3  df = df.replace({'RainToday': {'Yes': 1, 'No': 0}})

```

Listing 3.4: Encoding RainToday and RainTomorrow Variables

The 'Date' variable is dropped as it does not contribute to our analysis, and the rest of the categorical variables are encoded using the label encoder (Listing 3.5).

```

1  ### DROP THE COLUMN DATE
2  df.drop('Date', axis=1, inplace = True)
3
4  ### ENCODE THE CATEGORICAL VALUES
5  le = LabelEncoder() # creating instance of labelencoder
6  # Assigning numerical values to the columns
7  df['Location'] = le.fit_transform(df['Location'])
8  df['WindDir9am'] = le.fit_transform(df['WindDir9am'])
9  df['WindDir3pm'] = le.fit_transform(df['WindDir3pm'])
10 df['WindGustDir'] = le.fit_transform(df['WindGustDir'])

```

Listing 3.5: Encoding the categorical variables

As discussed earlier, the dataset contains potential outliers (data points that differ significantly from other observations). A suitable way to visualize them, is by plotting boxplots for the numerical variables.

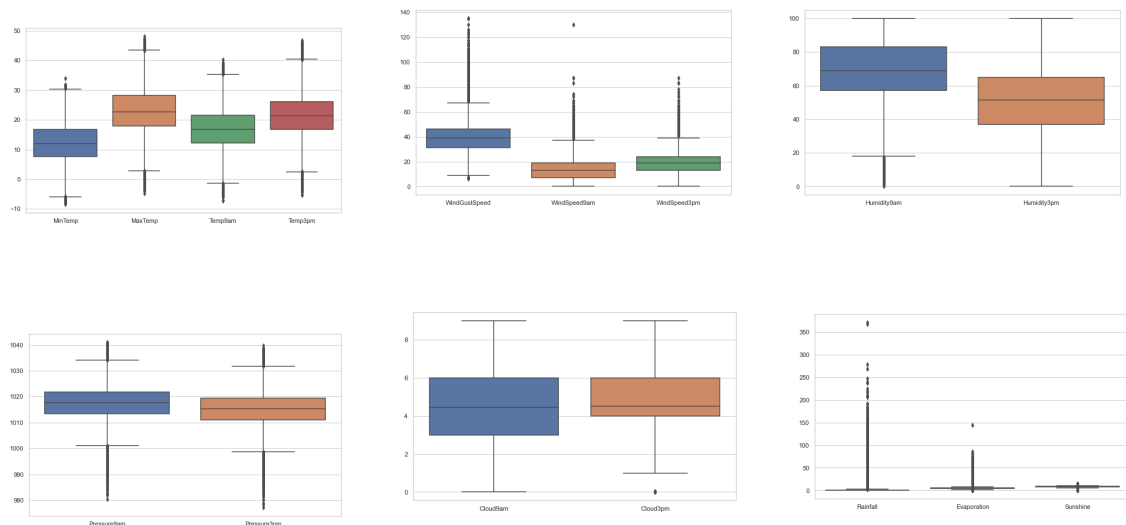


Figure 3.10: Boxplots of numerical variables before outlier treatment

Figure 3.10 confirms the presence of outliers. We can handle these outliers by using the z-score method. Z score tells how many standard deviations away a data point is from the mean. If the z score of a data point is more than 3, it indicates that the data point is quite different from the other data points. As shown in Listing 3.6, we keep all the observations with a z-score less than 3. We use the absolute value of the z-score as it can also be negative. We also print the shape of the dataset before and after outlier treatment.

```

1  ### HANDLE OUTLIERS AND PRINT SHAPE BEFORE AND AFTER
2
3  print('Shape of DataFrame Before Removing Outliers', df.shape )
4  df = df[(np.abs(stats.zscore(df)) < 3).all(axis=1)] # keep observations
   with z-score absolute value less than 3
5  print('Shape of DataFrame After Removing Outliers', df.shape )

```

Listing 3.6: Outlier treatment using z-score from scipy's stats module

We plot the boxplots again and we observe that outlier presence has been significantly improved (Figure 3.12).

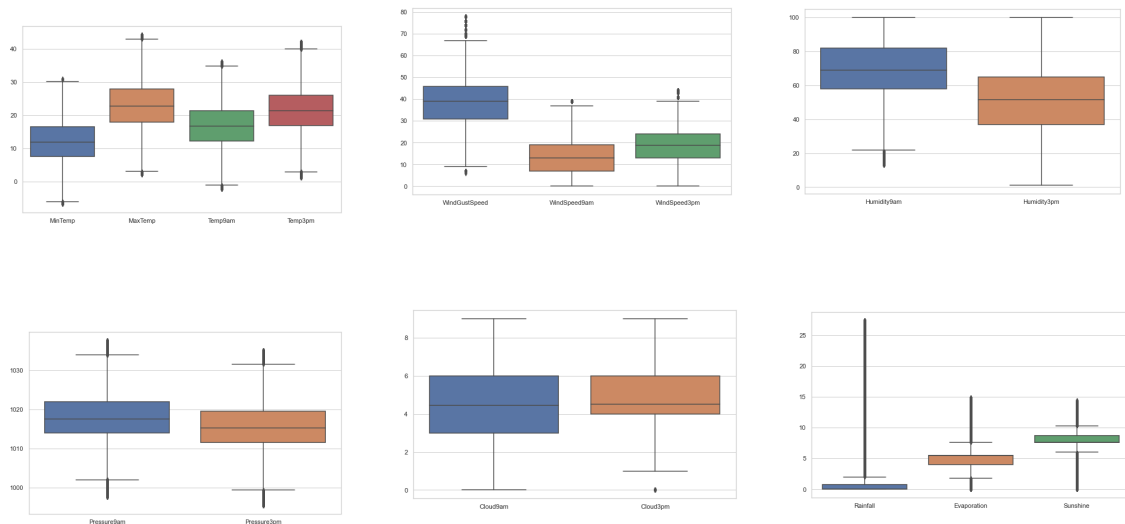


Figure 3.12: Boxplots of numerical variables after outlier treatment

The next step is data standardization (Listing 3.7). Standardization scales each input variable separately by subtracting the mean (called centering) and dividing by the standard deviation to shift the distribution to have a mean of zero and a standard deviation of one (Brownlee, 2020). The fit method is calculating the mean and variance of each of the features present in our data. The transform method is transforming all the features using the respective mean and variance. If we perform the fit method on our test data as well, we will get a new mean and variance for each feature, as well as allowing our model to learn about our test data (Khanna, 2020). This results to what is called data leakage (sharing the information between the test and training data sets).

```

1  ### SCALING THE DATA
2  # create and fit the scaler on training set
3  sc=StandardScaler()
4  X_train=sc.fit_transform(X_train)
5  X_test=sc.transform(X_test) # transform the test set

```

Listing 3.7: Data Standardization.

As shown earlier in Figure 3.6, the class variable is imbalanced. One approach to addressing imbalanced datasets is to oversample the minority class. The most widely used approach to synthesizing new examples is called the Synthetic Minority Oversampling TEchnique, or SMOTE (Brownlee, 2021b). To avoid data leakage, SMOTE is only applied on the training set (Listing 3.8). We print the number of samples of the class variable before and after applying SMOTE.

Before SMOTE: 0: 86697, 1: 22503  
 After SMOTE: 0: 86697, 1: 86697

```

1  %% OVERSAMPLE THE MINORITY CLASS AND PRINT THE NEW RATIO
2  # print number of samples before smote
3  counter = Counter(y_train)
4  print('Before SMOTE: ', counter)
5
6  # create the SMOTE oversampler
7  smt = SMOTE()
8
9  # fit oversampler on training set
10 X_smote, y_smote = smt.fit_resample(X_train, y_train)
11
12 # print number of samples after smote
13 counter = Counter(y_smote)
14 print('After SMOTE: ', counter)

```

Listing 3.8: Oversampling the minority class.

### 3.3 Model Training and Evaluation

To train and evaluate our model, we create the following functions:

The `fit_evaluate()` function (Listing 3.9) fits the model, predicts the labels, calculates training and testing time, prints the classification report and plots the confusion matrix.

```

1  # Function for fitting and evaluating the models
2  def fit_evaluate(model, x_train, y_train, x_test, y_test):
3      train_start = time.time() # model training starts
4      model.fit(x_train, y_train) # the model is fitted on the training set
5      train_stop = time.time() # model training stops
6      test_start=time.time() # testing starts
7      y_pred = model.predict(x_test) # model predicts the labels of the
      features in the test set
8      test_stop=time.time() # testing stops
9      print("Training time: ", train_stop - train_start) # print training
      time
10     print("Testing time: ", test_stop - test_start) # print testing time
11     print(classification_report(y_test, y_pred)) # print classification
      report
12
13     conf_matrix=confusion_matrix(y_test, y_pred) # calculate confusion
      matrix
14     ax = sns.heatmap(conf_matrix, annot=True, fmt = "g")
15     # plot the confusion matrix
16     ax.set_title('Confusion Matrix\n\n');
17     ax.set_xlabel('\nPredicted Values')
18     ax.set_ylabel('Actual Values ');
19
20
21     ## Display the visualization of the Confusion Matrix.
22     plt.show()

```

Listing 3.9: The `fit_evaluate()` function.

The `cross_validate()` function (Listing 3.10) calculates the 5-fold cross validation scores of the model, and prints their mean and standard deviation. This gives a more efficient use of

the data as every observation is used for both training and testing. Five fold cross validation was used instead of ten-fold to reduce the complexity.

```

1 def cross_validate(model, x_train, y_train):
2     scores=cross_val_score(model, x_train, y_train, cv=5) # calculate
    cross validation scores
3     print("Cross Validation Scores:\n")
4     print("Mean score: ", scores.mean()) # print the mean
5     print("Standard deviation: ", scores.std()) # print standard deviation
6     return scores

```

Listing 3.10: The cross\_validate() function.

The ROC\_AUC() function (Brownlee, 2021a) calculates and plots the ROC (Receiver Operating Characteristics) Curve of the model (Listing 3.11). ROC is a probability curve and AUC (Area under the curve) represents the degree or measure of separability. It tells how much the model is capable of distinguishing between classes. A perfect classifier will have a ROC AUC equal to 1, whereas a purely random classifier will have a ROC AUC equal to 0.5.

```

1 def ROC_AUC(model):
2     ns_probs = [0 for _ in range(len(y_test))] # generate a no skill
    prediction
3     lr_probs = model.predict_proba(X_test) # predict probabilities
4     lr_probs = lr_probs[:, 1] # keep probabilities for the positive
    outcome only
5     # calculate scores
6     ns_auc = roc_auc_score(y_test, ns_probs)
7     lr_auc = roc_auc_score(y_test, lr_probs)
8
9     # summarize scores
10    print('No Skill: ROC AUC=%.3f' % (ns_auc))
11    print('Model: ROC AUC=%.3f' % (lr_auc))
12
13    # calculate roc curves
14    ns_fpr, ns_tpr, _ = roc_curve(y_test, ns_probs)
15    lr_fpr, lr_tpr, _ = roc_curve(y_test, lr_probs)
16
17    # plot the roc curve for the model
18    plt.plot(ns_fpr, ns_tpr, linestyle='--', label='No Skill')
19    plt.plot(lr_fpr, lr_tpr, marker='.', label='Model')
20
21    # axis labels
22    plt.xlabel('False Positive Rate')
23    plt.ylabel('True Positive Rate')
24    # show the legend
25    plt.legend()
26    # show the plot
27    plt.show()

```

Listing 3.11: ROC\_AUC() function.

Next, we call the above functions to fit and evaluate our model. The algorithms used are Logistic Regression and Random Forest. Firstly, the models are fitted on the imbalanced data (Listing 3.12).

```

1 #%% ---- LOGISTIC REGRESSION -----
2
3 logreg = LogisticRegression(random_state=42,max_iter=700) # create the
    classifier
4 fit_evaluate(logreg, X_train, y_train, X_test, y_test) # fit and evaluate
    the model

```

```

5 scores_logreg=cross_validate(logreg,X_train, y_train) # print cross
  validation scores
6
7 ### ----- RANDOM FOREST -----
8
9 rfc=RandomForestClassifier() # create the classifier
10 fit_evaluate(rfc, X_train, y_train, X_test, y_test) # fit and evaluate the
  model
11 scores_rfc=cross_validate(rfc,X_train, y_train) # print cross validation
  scores

```

Listing 3.12: Fitting and evaluating the models with imbalanced data.

Next, the oversampled data is used as input (Listing 3.13)

```

1 ### ----- LOGISTIC REGRESSION ON OVERSAMPLED DATA -----
2
3 logreg_res=LogisticRegression(random_state=42,max_iter=700) # create the
  classifier
4 fit_evaluate(logreg_res, X_smote, y_smote, X_test, y_test) # fit and
  evaluate the model
5 log_res_scores=cross_validate(logreg_res,X_smote, y_smote) # print cross
  validation scores
6
7 ### ----- RANDOM FOREST ON OVERSAMPLED DATA -----
8 rfc_res=RandomForestClassifier() # create the classifier
9 fit_evaluate(rfc_res, X_smote, y_smote, X_test, y_test) # fit and evaluate
  the model
10 rfc_res_scores=cross_validate(rfc_res,X_smote, y_smote) # print cross
  validation scores

```

Listing 3.13: Fitting and evaluating the models with balanced data.

Multicollinearity occurs when features are highly correlated with one or more of the other features in the dataset. PCA (Principal Component Analysis) takes advantage of multicollinearity and combines the highly correlated variables into a set of uncorrelated variables. Therefore, PCA can effectively eliminate multicollinearity between features. We first need to find the number of principal components that capture an adequate percentage of information (Listing 3.14).

```

1 # Create and run a PCA
2 pca = PCA(n_components=None)
3 x_scaled = sc.fit_transform(X)
4 pca.fit(x_scaled) # fit the pca to scales x values (X is already scaled)
5
6 # Get the eigenvalues
7 print("Eigenvalues:")
8 print(pca.explained_variance_)
9 print()
10
11 # Get explained variances
12 print("Variances (Percentage):")
13 print(pca.explained_variance_ratio_ * 100)
14 print()
15
16 # Make the scree plot
17 plt.plot(np.cumsum(pca.explained_variance_ratio_ * 100))
18 plt.xlabel("Number of components (Dimensions)")
19 plt.ylabel("Explained variance (%)")

```

Listing 3.14: Running PCA to find optimal number of components.



Figure 3.13 shows the number of components and the percentage of their explained variance.




Figure 3.13: Number of components and the percentage of their explained variance.

We can see that we can capture around 90% of the information by using 11 principal components. We run PCA using `n_components=11` and create a new dataframe that contains these 11 components. We then split this dataset to training and test sets (Listing 3.15).

```

1 # run pca again with 11 principal components
2 x_scaled = sc.fit_transform(X)
3 pca = PCA(n_components=11)
4 X_pca = pca.fit_transform(x_scaled)
5
6 # Get the transformed dataset
7 X_pca = pd.DataFrame(X_pca)
8 print(X_pca.head())
9 print("\nSize: ")
10 print(X_pca.shape) # print the size of the dataframe
11 ###
12 # split PCA dataset to train and test sets
13 x_train_pca, x_test_pca, y_train_pca, y_test_pca = train_test_split(X_pca,
    y, test_size=0.20, shuffle=True, random_state=2)

```

Listing 3.15: Running PCA with 11 principal components and creating the new dataframe.

We then train our model using the training set that contains the principal components and evaluate it (Listing 3.16).

```

1 ### ----- LOGISTIC REGRESSION WITH PRINCIPAL COMPONENTS -----
2 logreg_pca = LogisticRegression(random_state=42,max_iter=700) # create the
    classifier
3 fit_evaluate(logreg_pca, x_train_pca, y_train_pca, x_test_pca, y_test_pca)
    # fit and evaluate the model
4 log_pca=cross_validate(logreg_pca,x_train_pca, y_train_pca) # print cross
    validation scores
5
6 ### ----- RANDOM FOREST WITH PRINCIPAL COMPONENTS -----
7 rfc_pca=RandomForestClassifier() # create the classifier
8 fit_evaluate(rfc_pca, x_train_pca, y_train_pca, x_test_pca, y_test_pca) #
    fit and evaluate the model
9 rfc_pca=cross_validate(rfc_pca,x_train_pca, y_train_pca) # print cross
    validation scores

```

Listing 3.16: Running PCA with 11 principal components.

As shown in Listing 3.17, to tune the hyperparameters and improve model performance, we run GridSearchCV for Logistic Regression and RandomizedSearchCV for Random Forest (GridSearchCV with Random Forest has a high computational time).

```

1  %% RANDOM FOREST WITH RANDOMIZED SEARCH CV
2  # Number of trees in random forest
3  n_estimators = [int(x) for x in np.linspace(start = 10, stop = 80, num =
    10)]
4  # Number of features to consider at every split
5  max_features = ['auto', 'sqrt']
6  # Maximum number of levels in tree
7  max_depth = [2,4]
8  # Minimum number of samples required to split a node
9  min_samples_split = [2, 5]
10 # Minimum number of samples required at each leaf node
11 min_samples_leaf = [1, 2]
12 # Method of selecting samples for training each tree
13 bootstrap = [True, False]
14
15 # Create the param grid
16 param_grid = {'n_estimators': n_estimators,
17               'max_features': max_features,
18               'max_depth': max_depth,
19               'min_samples_split': min_samples_split,
20               'min_samples_leaf': min_samples_leaf,
21               'bootstrap': bootstrap}
22
23 # create RandomizedSearchCV instance
24 rf_RandomGrid = RandomizedSearchCV(estimator = rfc, param_distributions =
    param_grid, cv = 10, verbose=2, n_jobs = 4)
25
26 rf_RandomGrid.fit(X_train, y_train) # fit to training sets
27 %%
28 print(rf_RandomGrid.best_params_) # print the best parameters

```

Listing 3.17: Random Forest Hyperparameter Tuning with RandomizedSearchCV.

The output parameters are: n\_estimators=72, min\_samples\_split=5, min\_samples\_leaf=1, max\_features='sqrt', max\_depth=4, bootstrap='True'.

Listing 3.18 shows the GridSearchCV performed for Logistic Regression. The parameters for the search are saved in a configuration file.

```

1  %% LOGISTIC REGRESSION GRIDSEARCH CV
2  # load contents from config file that contains the parameters
3  file="params.config"
4  contents=open(file).read()
5  parameters=eval(contents)
6  # create gridsearch cv instance
7  grid_search = GridSearchCV(estimator = logreg,
8                             param_grid = parameters,
9                             scoring = 'accuracy',
10                             cv = 5,
11                             verbose=0)
12
13 # fit the gridsearchcv to the training set
14 grid_search.fit(X_train, y_train)
15
16 %%
17

```

```

18 print('GridSearch CV best score : {:.4f}\n\n'.format(grid_search.
    best_score_))
19 print('Parameters that give the best results :', '\n\n', (grid_search.
    best_params_))
20 # print estimator that was chosen by the GridSearch
21 print('\n\nEstimator that was chosen by the search :', '\n\n', (grid_search
    .best_estimator_))

```

Listing 3.18: Logistic Regression Hyperparameter Tuning with GridSearchCV.

```

1 {
2 'penalty': ['l1', 'l2'],
3 'C': [1, 10, 100, 1000]
4 }

```

Listing 3.19: params.config.

The output parameters are: C=1, penalty: 'l2'. We then train both models using the hyperparameters acquired from the search.

```

1 ### ----- TUNED LOGISTIC REGRESSION -----
2 log_tuned=LogisticRegression(C=100, max_iter=700, random_state=42)
3 fit_evaluate(log_tuned, X_train, y_train, X_test, y_test)
4 scores_log_tuned=cross_validate(log_tuned, X_train, y_train)
5
6 ### TUNED RANDOM FOREST
7 rfc_tuned=RandomForestClassifier(n_estimators=72, min_samples_split=5,
    min_samples_leaf=1, max_features='sqrt', max_depth=4, bootstrap='True')
8 fit_evaluate(rfc_tuned, X_train, y_train, X_test, y_test)
9 scores_rfc_tuned=cross_validate(rfc_tuned, X_train, y_train)

```

Listing 3.20: Train models with tuned hyperparameters.

The results of our models are demonstrated below:

```

Training time: 0.17188715934753418
Testing time: 0.00040078163146972656
precision    recall  f1-score   support

      0       0.87       0.95       0.91       21684
      1       0.70       0.44       0.54       5617

 accuracy          0.85       27301
 macro avg          0.78       0.70       0.72       27301
weighted avg          0.83       0.85       0.83       27301

Cross Validation Scores:
Mean score: 0.8466941391941392
Standard deviation: 0.002360329110677981

```

Figure 3.14: Default Logistic Regression.

```

Training time: 17.34582495689392
Testing time: 0.49849510192871094
precision    recall  f1-score   support

      0       0.87       0.96       0.91       21684
      1       0.75       0.46       0.57       5617

 accuracy          0.86       27301
 macro avg          0.81       0.71       0.74       27301
weighted avg          0.85       0.86       0.84       27301

Cross Validation Scores:
Mean score: 0.8575274725274724
Standard deviation: 0.0022112303458226624

```

Figure 3.15: Default Random Forest.

```

Training time: 0.27912116050720215
Testing time: 0.001953125
precision    recall  f1-score   support

      0       0.92       0.79       0.85       21684
      1       0.48       0.75       0.58       5617

 accuracy          0.78       27301
 macro avg          0.70       0.77       0.72       27301
weighted avg          0.83       0.78       0.80       27301

Cross Validation Scores:
Mean score: 0.7763878858501222
Standard deviation: 0.0009480394266805599

```

Figure 3.16: Oversampled Logistic Regression.

```

Training time: 37.8157160282135
Testing time: 0.511972188949585
precision    recall  f1-score   support

      0       0.90       0.91       0.90       21684
      1       0.63       0.60       0.62       5617

 accuracy          0.85       27301
 macro avg          0.77       0.76       0.76       27301
weighted avg          0.84       0.85       0.85       27301

Cross Validation Scores:
Mean score: 0.9068365927144011
Standard deviation: 0.03833222482100837

```

Figure 3.17: Oversampled Random Forest.

```

Training time: 0.07933306694030762
Testing time: 0.0015387535095214844
precision    recall  f1-score   support

      0       0.86       0.95       0.90       21775
      1       0.67       0.40       0.50        5526

 accuracy          0.84       27301
  macro avg       0.77       0.67       0.70       27301
weighted avg       0.82       0.84       0.82       27301

Cross Validation Scores:

Mean score: 0.837087912087912
Standard deviation: 0.0009392587109859369

```

Figure 3.18: PCA Logistic Regression.

```

Training time: 31.8429958820343
Testing time: 0.5396690368652344
precision    recall  f1-score   support

      0       0.87       0.95       0.91       21775
      1       0.69       0.43       0.53        5526

 accuracy          0.85       27301
  macro avg       0.78       0.69       0.72       27301
weighted avg       0.83       0.85       0.83       27301

Cross Validation Scores:

Mean score: 0.8429853479853481
Standard deviation: 0.0015140107425286812

```

Figure 3.19: PCA Random Forest.

```

Training time: 0.21970272064208984
Testing time: 0.01021718978881836
precision    recall  f1-score   support

      0       0.87       0.95       0.91       21684
      1       0.70       0.44       0.54        5617

 accuracy          0.85       27301
  macro avg       0.78       0.70       0.72       27301
weighted avg       0.83       0.85       0.83       27301

Cross Validation Scores:

Mean score: 0.8466849816849816
Standard deviation: 0.002337515874594315

```

Figure 3.20: Tuned Logistic Regression.

```

Training time: 2.862377166748047
Testing time: 0.08210301399230957
precision    recall  f1-score   support

      0       0.84       0.98       0.91       21684
      1       0.81       0.27       0.40        5617

 accuracy          0.84       27301
  macro avg       0.83       0.63       0.65       27301
weighted avg       0.83       0.84       0.80       27301

Cross Validation Scores:

Mean score: 0.8358241758241759
Standard deviation: 0.002446375053488994

```

Figure 3.21: Tuned Random Forest.

### 3.4 Results

The results conclude the following:

- Oversampling produced a better F1 score compared to an imbalanced dataset in Logistic Regression, but significantly decreased the accuracy. For Random Forest, oversampling increased the F1 score without significantly reducing the accuracy.
- Compared to the default models, PCA improved the training time of Logistic Regression without reducing prediction capacity drastically, but increased computational time for Random Forest without producing better results.
- Tuning Logistic Regression produced the exact same results as the default model but decreased training time. Tuning Random Forest had decreased training time but did not bring significant gains that justified the loss of prediction capacity.
- All the cross validation scores are similar to the accuracies in the classification report. This means that all the models are stable. The small standard deviation of the scores also verifies this.

We come to the conclusion that the best Logistic Regression model is Tuned Logistic Regression and the best Random Forest model is the oversampled model. We now need to compare these two models. To do that, we have a look at their confusion matrices.

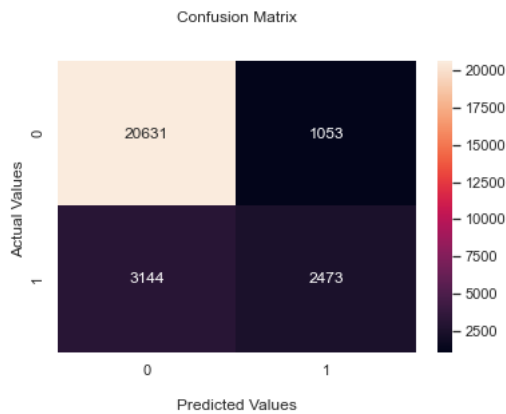


Figure 3.22: Logistic Regression Confusion Matrix.




Figure 3.23: Random Forest Confusion Matrix.

We can see that Logistic Regression predicts the negative class (No rainfall) better than Random Forest, while Random Forest predicts the positive class (Rainfall) more accurately. Here we need to think about the importance of correct classification. In this case, True Positives are more important than True Negatives as heavy rainfall in Australia is a bigger danger than an unexpected drought (Reuters, 2022). Considering this, the Random Forest model seems to be a better choice. To have a better look at the models' performance, we investigate the ROC curve and the AUC.




Figure 3.24: Logistic Regression ROC-AUC.




Figure 3.25: Random Forest ROC-AUC.

Since we can not tell a difference visually, we have a look at the metrics.

No Skill predictor: ROC AUC=0.500  
 Logistic Regression: ROC AUC=0.850  
 Random Forest: ROC AUC=0.874

Finally, we plot the cross validation scores of the two models using a box plot (Figure 3.26).




Figure 3.26: Cross validation scores.

The Logistic Regression model appears to be more stable, but Random Forest achieved significantly higher scores. However, we observe one score that is far off. Printing the cross validation scores, we see a lower score of 0.83 which is still a decent score.

Random forest scores: [0.83064679 0.90683122 0.93393696 0.93295654 0.93376204]

## Chapter 4

# Conclusions and Future Work

Based on the above performance results, the Random Forest algorithm trained on oversampled data outperformed any Logistic Regression model, with an accuracy of 85% and average cross validation score 90%. This score also outperforms the accuracies achieved by researchers that used Naive Bayes and Knn algorithms. Different approaches can be taken in the future in an attempt to achieve better results. Instead of feature extraction (PCA), feature selection can be implemented instead (e.g. SelectKBest). Also, the treatment of outliers can be revised. Some outliers can be legitimate observations and it is important to investigate their nature. For the evaluation of the models, it is also possible to check for overfitting by having a look at the accuracy scores over the training and test sets. Lastly, we can also test the ability of deep learning models to solving the problem of rainfall prediction.

# References

- AnalyticsVidhya (2017), 'Class imbalance: Handling imbalanced data using python'.  
**URL:** <https://www.analyticsvidhya.com/blog/2017/03/imbalanced-data-classification/>
- Australian Government Bureau of Meteorology (n.d.).  
**URL:** <http://www.bom.gov.au/>
- Brownlee, J. (2020), 'How to use standardscaler and minmaxscaler transforms in python'.  
**URL:** <https://machinelearningmastery.com/standardscaler-and-minmaxscaler-transforms-in-python/>
- Brownlee, J. (2021a), 'How to use roc curves and precision-recall curves for classification in python'.  
**URL:** <https://machinelearningmastery.com/roc-curves-and-precision-recall-curves-for-classification-in-python/>
- Brownlee, J. (2021b), 'Smote for imbalanced classification with python'.  
**URL:** <https://machinelearningmastery.com/smote-oversampling-for-imbalanced-classification/>
- Gupta, D. and Ghose, U. (2015), A comparative study of classification algorithms for forecasting rainfall, in '2015 4th International Conference on Reliability, Infocom Technologies and Optimization (ICRITO) (Trends and Future Directions)', pp. 1–6.
- Khanna, C. (2020), 'What and why behind `fit_transform()` vs `transform()` in `scikit-learn`!',  
**URL:** <https://towardsdatascience.com/what-and-why-behind-fit-transform-vs-transform-in-scikit-learn-78f915c>
- MachineLearningMastery (2021), 'Failure of classification accuracy for imbalanced class distributions'.  
**URL:** <https://machinelearningmastery.com/failure-of-accuracy-for-imbalanced-class-distributions/>
- Nicholls, N. (2001), 'Atmospheric and climatic hazards: Improved monitoring and prediction for disaster mitigation. natural hazards'.
- Reuters (2022), 'Sydney faces more rain as death toll from australian floods rises'.  
**URL:** <https://edition.cnn.com/2022/03/06/asia/sydney-australia-flood-intl-hnk/index.html>
- Sharma, K. K., Verdon-Kidd, D. C. and Magee, A. D. (2021), 'A decision tree approach to identify predictors of extreme rainfall events – a case study for the fiji islands', *Weather and Climate Extremes* **34**, 100405.  
**URL:** <https://www.sciencedirect.com/science/article/pii/S221209472100092X>
- Sharmila (2020), 'Mechanisms of multiyear variations of northern australia wet-season rainfall'.



Wand (n.d.).

**URL:** <https://wand.net.nz/pubs/22/html/node34.html>

WHO (2003), 'Climate change and human health: Risks and re- sponses.'